

LINUX系统高级调试和优化

-- 信号、异常处理和应用程序崩溃

张银奎

2016/8/2(2017/12 UPDATE)

1

信号和信号处理



崩溃

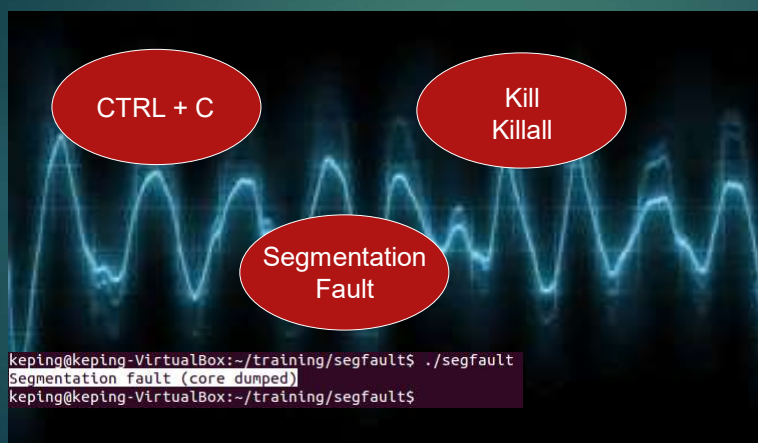


错误报告
(Apport)

2

无处不在的SIGNAL

Linux系统每一秒中都有大量的信号在产生、传递、接收和处理，不管你有没有意识到它的存在




从本页开始的信号部分，由王科平为第一届庐山研习班而编写

3

什么是SIGNAL

- 在软件层次上是对中断机制的一种模拟
- 一种进程间的异步通信机制
- 内核也可以因为内部事件而给进程发送信号，通知进程发生了某个事件 (如著名的segmentation fault)
- 从UNIX继承过来，Linux进行了扩展和改进



4

Inside SIGNAL

`include/linux/sched.h`

```

Struct task_struct{
    ...
    /* signal handlers */
    struct signal_struct *signal; // Pointer to the process's signal descriptor
    struct sighand_struct *sighand; // Pointer to the process's signal handler descriptor

    sigset_t blocked, real_blocked; // Mask of blocked signals
    sigset_t saved_sigmask; // restored if set_restore_sigmask() was used
    struct sigpending pending; // the private pending signals

    unsigned long sas_ss_sp; // alternative signal handler stack
    size_t sas_ss_size; // size of alternative signal handler stack
    int (*notifier)(void *priv); // Pointer to a function used by a device driver to block some signals of the process
    void *notifier_data; // data that might be used by the notifier function
    sigset_t *notifier_mask; // Bit mask of signals blocked by a device driver through a notifier function
    ...
}

```

5


Inside SIGNAL

场景1：跨进程通讯

Step 1: 进程A通过signal系统调用注册一个signal handler
(optional, 如未注册, 将有缺省处理)

Step 2: 进程B通过kill系统调用触发signal

Step 3: 内核太sys_kill处理进程B的调用请求, 触发进程A的signal handler



6

Inside SIGNAL

场景2：进程运行中的异常触发signal

- Step 1: 进程通过signal系统调用注册一个signal handler
 (optional, 如未注册, 将有缺省处理)
- Step 2: 进程运行过程中触发系统调用 (比如memory处理)
- Step 3: 内核态处理系统调用过程中出现异常, 触发进程A的signal handler
 最常见的就是segmentation fault

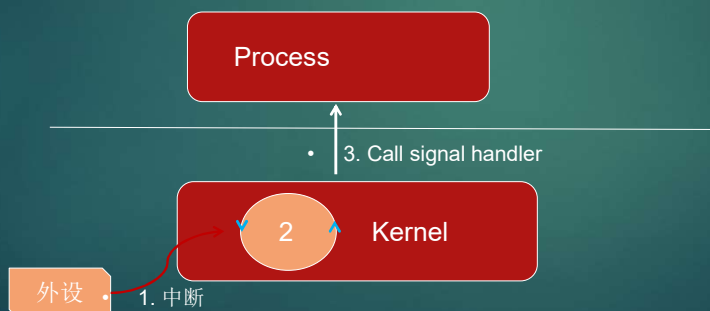


7

Inside SIGNAL

场景3：外部中断异常触发signal

- Step 1: 进程通过signal系统调用注册一个signal handler
 (optional, 如未注册, 将有缺省处理)
- Step 2: 内核收到外设中断 (比如键盘 CTRL+C)
- Step 3: 内核态中断处理过程中触发进程A的signal handler
 1). 处理系统调用过程中出现异常
 2). CTRL + C, CTRL + Z



8

Legacy (Unix) vs Linux

- ▶ Signal 1~32 从Unix继承而来
- ▶ 33~64 是Linux才有的
- ▶ Legacy Signals:
 - ▶ 不可靠
 - ▶ 无法传递数据
 - ▶ 通常用kill触发
- ▶ Linux Signals
 - ▶ 可靠
 - ▶ 可以传递数据
 - ▶ 通常用sigqueue触发

9

9

Legacy (UNIX) Signals

```

#define SIGHUP      1      /* Hangup (POSIX).  */
#define SIGINT      2      /* Interrupt (ANSI). */
#define SIGQUIT     3      /* Quit (POSIX).   */
#define SIGILL      4      /* Illegal instruction (ANSI). */
#define SIGTRAP     5      /* Trace trap (POSIX). */
#define SIGIOT      6      /* IOT trap (4.2 BSD). */
#define SIGABRT     SIGIOT /* Abort (ANSI).   */
#define SIGEMT      7
#define SIGFPE      8      /* Floating-point exception (ANSI). */
#define SIGKILL     9      /* Kill, unblockable (POSIX). */
#define SIGBUS      10     /* BUS error (4.2 BSD). */
#define SIGSEGV     11     /* Segmentation violation (ANSI). */
#define SIGSYS      12
#define SIGPIPE     13     /* Broken pipe (POSIX). */
#define SIGALRM     14     /* Alarm clock (POSIX). */
#define SIGTERM     15     /* Termination (ANSI). */
#define SIGUSR1     16     /* User-defined signal 1 (POSIX). */

```

10

10

Legacy (UNIX) Signals

```

#define SIGUSR2      17      /* User-defined signal 2 (POSIX). */
#define SIGCHLD      18      /* Child status has changed (POSIX). */
#define SIGCLD       SIGCHLD /* Same as SIGCHLD (System V). */
#define SIGPWR       9       /* Power failure restart (System V). */
#define SIGWINCH     20      /* Window size change (4.3 BSD, Sun). */
#define SIGURG       21      /* Urgent condition on socket (4.2 BSD). */
#define SIGIO        22      /* I/O now possible (4.2 BSD). */
#define SIGPOLL       SIGIO  /* Pollable event occurred (System V). */
#define SIGSTOP      23      /* Stop, unblockable (POSIX). */
#define SIGTSTP      24      /* Keyboard stop (POSIX). */
#define SIGCONT       25      /* Continue (POSIX). */
#define SIGTTIN       26      /* Background read from tty (POSIX). */
#define SIGTTOU       27      /* Background write to tty (POSIX). */
#define SIGVTALRM     28      /* Virtual alarm clock (4.2 BSD). */
#define SIGPROF       29      /* Profiling alarm clock (4.2 BSD). */
#define SIGXCPU       30      /* CPU limit exceeded (4.2 BSD). */
#define SIGXFSZ       31      /* File size limit exceeded (4.2 BSD). */

```

11

用户定义信号

► 使用USR1和USR2

```

// include/linux/signal.h
/* Test if 'sig' is valid signal. Use this instead of testing _NSIG directly */
static inline int valid_signal(unsigned long sig)
{
    return sig <= _NSIG ? 1 : 0;
}

```

```

int send_sig_info(int sig, struct siginfo *info, struct task_struct *p)
{
    /*
     * Make sure legacy kernel users don't send in bad values
     * (normal paths check this in check_kill_permission).
     */
    if (!valid_signal(sig))
        return -EINVAL;

    return do_send_sig_info(sig, info, p, false);
}

```

```
// arch/x86/include/asm/signal.h
```

```
#define _NSIG      64
```

12

相关的系统调用

- ▶ Signal
- ▶ Kill
- ▶ Sigaction
- ▶ Sigqueue
- ▶ . . .

13

13

进阶：屏蔽信号

```
sigset_t mask;

/* 创建一个 mask，加入要屏蔽的信号*/
sigemptyset(&mask);
sigaddset(&mask, SIGINT); //add "Ctrl +C" signal

/* 通过系统调用，设置新的mask */
sigprocmask(SIG_BLOCK,&mask,NULL);

/* 解除屏蔽 */
sigprocmask(SIG_UNBLOCK,&mask,NULL);
```

14

14

进阶：跳过异常继续运行

- ▶ 通过signal + setjmp / longjmp使进程在收到导致异常退出的的signal时候，跳过异常，继续运行
- ▶ 其中setjmp用于保存执行状态，longjmp用于跳回setjmp保存的状态
- ▶ 最佳的应用模式是：处理异常情况，使进程优雅地退出

15

15

理解setjmp/longjmp

```
#include < setjmp.h >
```

```
main()
```

```
{
```

```
    jmp_buf env;
```

```
    int i;
```

```
    i = setjmp(env);
```

```
    printf("i = %d\n", i);
```

```
    if (i != 0) exit(0);
```

```
    longjmp(env, 2);
```

```
    printf("Does this line get printed?\n");
```

```
}
```

- ▶ 编译执行

- ▶ i = 0

- ▶ i = 2

- ▶ 第一次执行setjmp返回0
- ▶ 调用longjmp时，CPU再次回去执行setjmp，并返回longjmp指定的返回值2

16

DEMO

- ▶ 体验signal
 - ▶ 屏蔽CTRL + C
 - ▶ SEGV (segmentation fault初体验)
 - ▶ 跳过segmentation fault，继续运行

代码在 **home/keping/training/signal** 目录下
Signal.c的第9到12行定义了四个宏：

```
HANDLE_SIGNAL  
DO_SEGFAULT  
JUMP_SEGFAULT  
HANDLE_CTRL_C
```

可以通过注释掉或启用这几个宏来开关某些功能

17

17

Handle SIGSEGV (段错误)

- 开关HANDLE_SIGNAL，体验并理解不同的结果

18

注：在这个实验中保持

- 1.关闭JUMP_SEGFAULT和HANDLE_CTRL_C（注释掉第10，12行）
- 2.一直保持开启DO_SEGFAULT(第10行)

18

跳过段错误

- 开启HANDLE_SIGNAL, DO_SEGFAULT和JUMP_SEGFAULT, 关闭HANDLE_CTRL_C, 体验通过setjump / longjmp处理段错误的效果

```
ge@gewubox:~/work/signal$ ./signal
Process started
Bang! going to segment fault!
Child: received signal 11: SEGMENT FAULT
Recovered from segment fault!!
Process ended
```

源代码位于setjump/longjmp.c
X64的核心实现在sysdep/x86_64/__longjmp.S

19

屏蔽CTRL + C

- 开关HANDLE_SIGNAL和HANDLE_CTRL_C, 重新编译并运行 ./signal
- Signal运行起来后, 按CTRL+C

```
ge@gewubox:~/work/signal$ ./ctrlc
Process started
SIGINT is blocked for 10 seconds, CTRL+C takes no effect now
SIGINT is unblocked, CTRL+C works now
Sleep for 10 seconds, or use kill or CTRL+C to exit
^C
ge@gewubox:~/work/signal$
```

注: 在这个实验中关闭DO_SEGFAULT 和JUMP_SEGFAULT (注释掉第9, 10行)

20

小结

原理：

- 在软件层次上是对中断机制的一种模拟
- 一种进程间的异步通信机制
- 内核也可以因为内部事件而给进程发送信号，通知进程发生了某个事件 (如著名的 segmentation fault)
- 从UNIX继承过来，Linux进行了扩展和改进

基本操作：

1. 注册signal处理函数：signal, sigaction
2. 触发signal：kill, sigqueue

应用：

1. 进程间通讯
2. 通过自定义的signal处理函数，更优雅的处理异常处理

21

信号和信号处理



崩溃

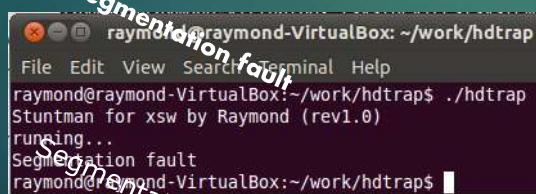


错误报告
(Apport)

22

背景

- ▶ 重要软件的后台服务随机崩溃
- ▶ 著名的段错误
 - ▶ 在控制台窗口运行
 - ▶ 留下一句话Segmentation fault
- ▶ “替身”小程序，模拟故障



```

raymond@raymond-VirtualBox: ~/work/hdtrap
File Edit View Search Terminal Help
raymond@raymond-VirtualBox:~/work/hdtrap$ ./hdtrap
Stuntman for xsw by Raymond (rev1.0)
running...
Segmentation fault
raymond@raymond-VirtualBox:~/work/hdtrap$

```

23

demo

感受段错误

24

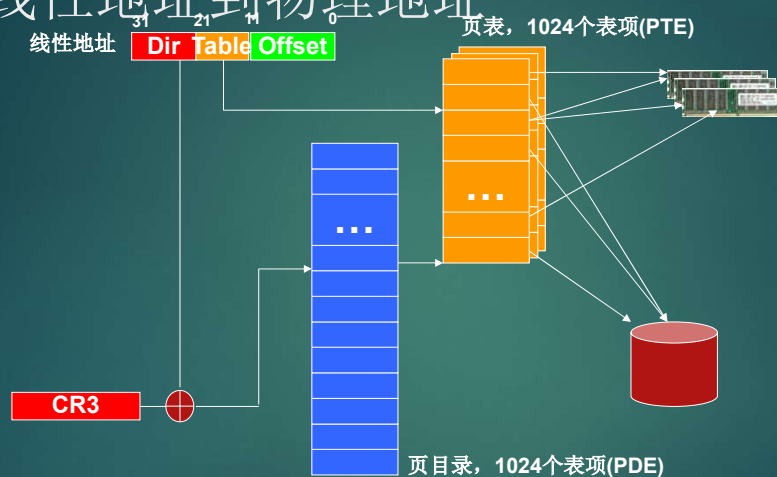
保护模式



25

25

从线性地址到物理地址




26

- ▶ 上面是以4KB内存页(未启用PAE)为例
- ▶ $1024\text{PDE} \times 1024\text{PTE} = 2^{20}$ 页

26

特权和秩序



DPL – Descriptor Privilege Level



U/S – User(1)/Supervisor(0)



- ▶ 正在执行的代码所在代码段的特权级别就是该代码的特权级别
- ▶ 当一段代码调用位于其它段的函数或访问其它段的数据时，CPU会检查发起访问者是否有足够的权限。如果没有通过检查，则产生异常

27

空指针

- ▶ NULL Pointer
- ▶ In computing, a null pointer has a value reserved for indicating that the pointer does not refer to a valid object.

I call it my billion-dollar mistake. It was the invention of the null reference in 1965. At that time, I was designing the first comprehensive type system for references in an object oriented language (ALGOL W). My goal was to ensure that all use of references should be absolutely safe, with checking performed automatically by the compiler. But I couldn't resist the temptation to put in a null reference, simply because it was so easy to implement. This has led to innumerable errors, vulnerabilities, and system crashes, which have probably caused a billion dollars of pain and damage in the last forty years.



Charles Antony Richard Hoare

28

违反规则

- ▶ 异常
- ▶ Linux
 - ▶ 段错误
 - ▶ 名字比实际范围小
- ▶ Windows
 - ▶ 访问违例
 - ▶ Access violation
 - ▶ 0xC0000005



29

29

dmesg

```
[25402.126121] hdtrap[1960]: segfault at 0 ip 080484b8 sp bfee0c28 error 6 in
hdtrap[8048000+1000]
```

- ▶ IP地址为0x080484b8的指令访问地址0导致了段错误，当时的栈指针是0xbfee0c28，错误码是6（用户态的写操作，参见下文），指令所属的内存块的基地址为8048000，长度为0x1000

30

30

```
// kernel-2.6.35\arch\x86\mm\fault.c
/*
 * Print out info about fatal segfaults, if the show_unhandled_signals
 * sysctl is set:
 */
static inline void
show_signal_msg(struct pt_regs *regs, unsigned long error_code,
                unsigned long address, struct task_struct *tsk)
{
    if (!unhandled_signal(tsk, SIGSEGV))
        return;
    if (!printk_ratelimit())
        return;
    printk("%s%s[%d]: segfault at %lx ip %p sp %p error %lx",
          task_pid_nr(tsk) > 1 ? KERN_INFO : KERN_EMERG,
          tsk->comm, task_pid_nr(tsk), address,
          (void *)regs->ip, (void *)regs->sp, error_code);

    print_vma_addr(KERN_CONT " in ", regs->ip);
    printk(KERN_CONT "\n");
}
```

31

31

错误码

```
/*
 * Page fault error code bits:
 */
* bit 0 == 0: no page found    1: protection fault
* bit 1 == 0: read access      1: write access
* bit 2 == 0: kernel-mode access 1: user-mode access
* bit 3 ==                    1: use of reserved bit detected
* bit 4 ==                    1: fault was an instruction fetch
*/
enum x86_pf_error_code {

    PF_PROT =    1 << 0,
    PF_WRITE =   1 << 1,
    PF_USER =    1 << 2,
    PF_RSVD =    1 << 3,
    PF_INSTR =   1 << 4,
};
```

32

32

Map文件

► 产生Map文件

- gcc -o hdtrap hdtrap.c md5.c -Wl,-Map=hdtrap.map
- 注意W后是小写L，代表链接选项

```

hdtrap.map
Archive member included because of file (symbol)
//usr/lib/libc_nonshared.a(elf-init.oS)
/usr/lib/gcc/i686-linux-gnu/4.4.5/../../../../lib/crt1.o (__libc_csu_fini)

Discarded input sections

.note.GNU-stack           0x0000000000000000      0x0 /usr/lib/gcc/i686-linux-gnu/4.4.5/../../../../lib/crt1.o
.gnu_debuglink             0x0000000000000000      0xc /usr/lib/gcc/i686-linux-gnu/4.4.5/../../../../lib/crt1.o
.note.GNU-stack           0x0000000000000000      0x0 /usr/lib/gcc/i686-linux-gnu/4.4.5/../../../../lib/crti.o
.gnu_debuglink             0x0000000000000000      0xc /usr/lib/gcc/i686-linux-gnu/4.4.5/../../../../lib/crti.o
.note.GNU-stack           0x0000000000000000      0x0 /usr/lib/gcc/i686-linux-gnu/4.4.5/crtbegin.o
.note.GNU-stack           0x0000000000000000      0x0 /tmp/ccAUceXL.o
.note.GNU-stack           0x0000000000000000      0x0 /tmp/ccp80jT3.o
.group                     0x0000000000000000      0x8 /usr/lib/libc_nonshared.a(elf-init.oS)

```

33

查找

- [27566.531498] hdtrap[2504]: segfault at 0 ip 080484b8 sp bff209d8 error 6 in hdtrap[8048000+1000]

```

.text           0x000000008048394      0x0 /usr/lib/gcc/i686-linux-gnu/4.4.5/../../../../lib/crti.o
*fill*         0x000000008048394      0xc 90909090
.text          0x0000000080483a0      0x83 /usr/lib/gcc/i686-linux-gnu/4.4.5/crtbegin.o
*fill*         0x000000008048423      0x1 90909090
.text          0x000000008048424      0x86 /tmp/ccAUceXL.o
               0x000000008048424      get_file_id
               0x00000000804846e      main
*fill*         0x0000000080484aa      0x2 90909090
.text          0x0000000080484ac      0x15 /tmp/ccp80jT3.o
               0x0000000080484ac      calc_md5
*fill*         0x0000000080484c1      0xf 90909090
.text          0x0000000080484d0      0x6a /usr/lib/libc_nonshared.a(elf-init.oS)
               0x0000000080484d0      __libc_csu_fini
               0x0000000080484e0      __libc_csu_init

```

34

有请GDB

► GNU Debugger

- 1986年, Richard Stallman创建
- 1900-1993, John Gilmore维护
- 目前在GDB Steering Committee
- 最新版本7.3 (2011.7.26)
- 支持很多种CPU architecture
- A29K, ARC, ETRAX CRIS, D10V, D30V, FR-30, FR-V, Intel i960, M32R, 68HC11, Motorola 88000, MCORE, MN10200, MN10300, NS32K, Stormy16, V850, Z8000 and many more
- 默认为命令行界面,有很多GUI的前端(Add-on)



35

35

常用调试命令对照表

WinDBG命令	GDB命令	功能
bp	break 或 b	设置软件断点
ba	watch	设置硬件断点
k	backtrace 或 bt	显示函数调用序列 (栈回溯)
g	continue 或 c	恢复执行
p/t	next/step 或 n/s	单步跟踪
d	x	观察内存
dv	info locals	观察局部变量
dt	pt	观察数据类型 (结构)
gu	finish	执行到函数返回
.frame	frame	切换当前栈帧
lm	i shared	列模块

36

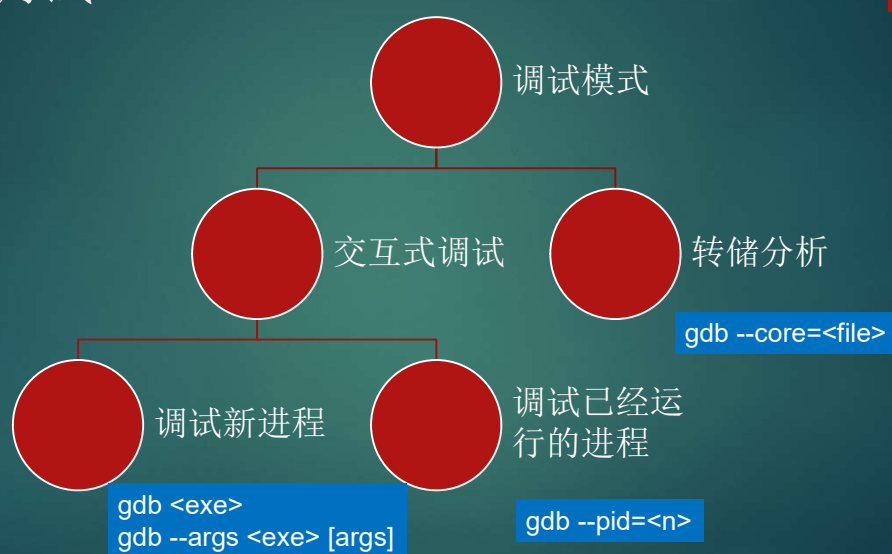
更多命令

- ▶ 命令行
 - ▶ run xxx
 - ▶ set args xx xx
 - ▶ show args
- ▶ 观察类型
 - ▶ ptype
 - ▶ whatis
 - ▶ print v@10
- ▶ 管理断点
 - ▶ info/disable/delete break
- ▶ 源代码
 - ▶ list 3,8
 - ▶ info line/source/sources
- ▶ show conv

37

37

开始调试



38

38

或者

- ▶ gdb
- ▶ file <exe>
- ▶ run [args]

```

raymond@ubuntu: ~/work/hdtrap
File Edit View Search Terminal Help
raymond@ubuntu:~$ cd work
raymond@ubuntu:~/work$ cd hdtrap
raymond@ubuntu:~/work/hdtrap$ gdb
GNU gdb (GDB) 7.2-ubuntu
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "i686-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
(gdb) file ./hdtrap
Reading symbols from /home/raymond/work/hdtrap/hdtrap...done.
(gdb) run -args
Starting program: /home/raymond/work/hdtrap/hdtrap -args
Stuntman for xsw by Raymond (rev1.0)
running...

Program received signal SIGSEGV, Segmentation fault.
0x080484b8 in calc_md5 (data=0x8048590 "testing data-xxxxxxx", nLen=20,
md5=0x0) at md5.c:7
7       md5[0] = A;
(gdb)

```

39

39

案发函数

```

Program received signal SIGSEGV, Segmentation fault.
0x080484b8 in calc_md5 (data=0x8048590 "testing data-xxxxxxx", nLen=20,
md5=0x0) at md5.c:7
7       md5[0] = A;

```

40

```

int calc_md5(char * data,
             int nLen,
             unsigned int md5[4])
{
    int A, B, C, D;

    // calc md5 for data array now

    // assign the result now
    md5[0] = A;
    // ...

    return 0;
}

```



md5等于0，明显访问空指针

40

追查调用者

```
(gdb) bt
#0  0x080484b8 in calc_md5 (data=0x8048590 "testing data-xxxxxxx", nLen=20,
md5=0x0) at md5.c:7
#1  0x08048468 in get_file_id (filename=0x80485d8 "filea", fileid=0xbffff3d0)
at hdtrap.c:24
#2  0x080484a3 in main (argc=1, argv=0xbffff494) at hdtrap.c:35
```

```
int get_file_id(char * filename, unsigned int * fileid)
{
    SECTION_HEAD section;
    section.data = "testing data-xxxxxxx";
    section.length = strlen(section.data);
    return calc_md5(section.data,
        section.length, fileid);
}
```

也是来自参数

41

继续上溯

```
(gdb) frame 2
#2  0x080484a3 in main (argc=1, argv=0xbffff494) at hdtrap.c:35
35      get_file_id("filea", fileid);
(gdb) info locals
fileid = {1436037, 1174400, 134513899, 2658292}
(gdb) p &fileid
$3 = (unsigned int (*)[4]) 0xbffff3d0
```

- ▶ 在main这一级，传递下去的是局部数组变量fileid，显然不为空
- ▶ 传了两次就出问题了吗？

42

主函数

43

```

28 int main(int argc, char * argv[])
29 {
30     unsigned int fileid[4];
31
32     printf("Stuntman for xsw by Raymond (rev1.0)\n");
33     printf("running...\n");
34
35     get_file_id("filea", fileid);
36
37     return getchar();
38 }

```

43

审查传递过程

44

```

0x08048448 <+36>:  mov    -0x10(%ebp),%eax
0x0804844b <+39>:  mov    -0xc(%ebp),%edx
0x0804844e <+42>:  mov    -0x14(%ebp),%ecx
0x08048451 <+45>:  mov    0xc(%ebp),%ebx
0x08048454 <+48>:  mov    %ebx,0xc(%esp)
0x08048458 <+52>:  mov    %eax,0x4(%esp)
0x0804845c <+56>:  mov    %edx,0x8(%esp)
0x08048460 <+60>:  mov    %ecx,(%esp)
0x08048463 <+63>:  call   0x80484ac <calc_md5>
=> 0x08048468 <+68>:  add    $0x24,%esp
0x0804846b <+71>:  pop    %ebx

```

44

X86汇编语言

INTEL语法

- 先是目标，然后是源，也就是从右向左赋值
- Windows上流行

AT&T语法

- 先是源，然后是目标，也就是从左向右赋值
- Unix和Linux上流行

45

set disassembly-flavor

```
set disassembly-flavor att
set disassembly-flavor intel
show disassembly-flavor
```

http://visualgdb.com/gdbreference/commands/set_disassembly-flavor

46

关键数据结构

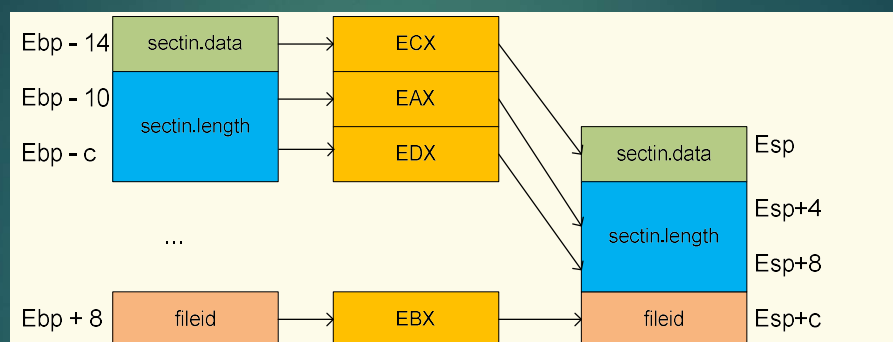
```
(gdb) pt section
type = struct {
  char *data;
  ULONGLONG length;
}
```

```
(gdb) p sizeof(section.length)
$3 = 8
```

```
(gdb) p sizeof(section)
$8 = 12
```

47

准备参数

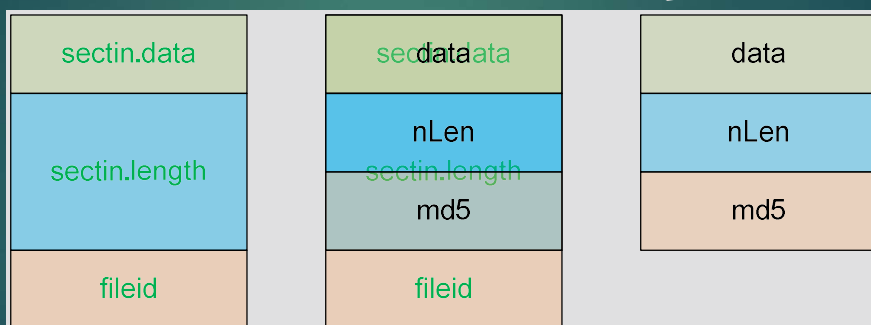


48

再看崩溃现场

```
0x080484b2 <+6>: mov  -0x4(%ebp),%edx
0x080484b5 <+9>: mov  0x10(%ebp),%eax
0x080484b8 <+12>: mov  %edx,(%eax)
```

49



父函数准备的参数

子函数使用的参数

49

函数原型 - 契约

```
int calc_md5(char * data,
             int nLen,
             unsigned int md5[4])
```

第二个参数是int,
可父函数为什么
非要传递64位呢?

50

这样的错误让人囧



50

要点

- ▶ 段错误
 - ▶ 违反了系统的保护规则
- ▶ dmesg查看内核记录的摘要
- ▶ 使用map文件初步定位
- ▶ 使用GDB
- ▶ C标准，对于未声明函数使用默认约定，可能“严重误会”
- ▶ 声明原型或者严格包含头文件

51

51

JIT调试Linux程序

- ▶ 注册信号处理器
- ▶ 触发时启动调试器进程

52

```
sigact.sa_handler = signal_jit_handler;

sigaddset(&sigact.sa_mask, SIGSEGV);
sigaction(SIGSEGV, &sigact, (struct sigaction *)NULL);

char gdb[160];

sprintf(gdb, "gdb --pid %d", getpid());
system(gdb);
```

52

```

static void crash_handler(int sig)
{
    int status=0;
    int pid;
    char * gdb_array[]={ "gdb", "", NULL};
    char pid_str[40];

    sprintf(pid_str, "--pid=%d%c", getpid(), '\0');
    gdb_array[1]= pid_str;

    pid= fork();

    if (pid < 0) /* error */
        abort();
    else if (pid) /* parent */
    {
        while (1)
            sleep(6000); /* Give GDB time to attach */
        // _exit(1); /* you can skip this line by telling gdb to "return" */
    }
    else /* child */
        execvp("gdb", gdb_array);
}

```

53

注册信号处理器

```

void register_jit_gdb()
{
    signal(SIGQUIT, crash_handler); /* Normally got from Ctrl-\ */
    signal(SIGILL, crash_handler);
    signal(SIGTRAP, crash_handler);
    signal(SIGABRT, crash_handler);
    signal(SIGFPE, crash_handler);
    signal(SIGBUS, crash_handler);
    signal(SIGSEGV, crash_handler); /* This is the most common crash */
    signal(SIGSYS, crash_handler);
}

```

54

失败

```

ge@gewubox: ~/work/hdtrap
hdtrap hdtrap.map hdtrapr Makefile md5.o
ge@gewubox:~/work/hdtrap$ gedit hdtrap.c &
[1] 3120
ge@gewubox:~/work/hdtrap$ ./hdtrap -jit
Stuntman for xsw by Raymond (rev1.0)
running...
jit debug handler registered
GNU gdb (Ubuntu/Linaro 7.4-2012.04-0ubuntu2.1) 7.4-2012.04
Copyright (C) 2012 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "i686-linux-gnu".
For bug reporting instructions, please see:
<http://bugs.launchpad.net/gdb-linaro/>.
Attaching to process 3136
Could not attach to process. If your uid matches the uid of the target
process, check the setting of /proc/sys/kernel/yama/ptrace_scope, or try
again as the root user. For more details, see /etc/sysctl.d/10-ptrace.conf
ptrace: Operation not permitted.
(gdb) bt
No stack.

```

55

ptrace_scope

- ▶ Ubuntu 10.10引入的安全机制
- ▶ 默认值为1，防止一个进程分析和修改另一个进程，除非后者是前者的子进程

TEE(1)	User Commands	TEE(1)
NAME	tee - read from standard input and write to standard output and files	
SYNOPSIS	tee [OPTION]... [FILE]...	
DESCRIPTION	Copy standard input to each FILE, and also to standard output.	
	-a, --append append to the given FILES, do not overwrite	

echo "0" | sudo tee /proc/sys/kernel/yama/ptrace_scope

56

```

root@gewubox: /home/ge/work/hdtrap
root@gewubox:/home/ge/work/hdtrap# ./hdtrap
Stuntman for xsw by Raymond (rev1.0)
running...
Segmentation fault (core dumped)
root@gewubox:/home/ge/work/hdtrap# ./hdtrap -jit
Stuntman for xsw by Raymond (rev1.0)
running...
jit debug handler registered
GNU gdb (Ubuntu/Linaro 7.4-2012.04-0ubuntu2.1) 7.4-2012.04
Copyright (C) 2012 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "i686-linux-gnu".
For bug reporting instructions, please see:
<http://bugs.launchpad.net/gdb-linaro/>.
Attaching to process 3288
Reading symbols from /home/ge/work/hdtrap/hdtrap...done.
Reading symbols from /lib/i386-linux-gnu/libm.so.6...Reading symbols from /usr/l
ib/debug/lib/i386-linux-gnu/libm-2.15.so...done.
done.
Loaded symbols for /lib/i386-linux-gnu/libm.so.6
Reading symbols from /lib/i386-linux-gnu/libpthread.so.0...Reading symbols from

```

57

```

done.
Loaded symbols for /lib/i386-linux-gnu/libpthread.so.0
Reading symbols from /lib/i386-linux-gnu/libc.so.6...Reading symbols from /usr/l
ib/debug/lib/i386-linux-gnu/libc-2.15.so...done.
done.
Loaded symbols for /lib/i386-linux-gnu/libc.so.6
Reading symbols from /lib/ld-linux.so.2...Reading symbols from /usr/lib/debug/li
b/i386-linux-gnu/ld-2.15.so...done.
done.
Loaded symbols for /lib/ld-linux.so.2
0xb7772424 in __kernel_vsyscall ()
(gdb) bt
#0  0xb7772424 in __kernel_vsyscall ()
#1  0xb76270e0 in __nanosleep_nocancel ()
    at ../sysdeps/unix/syscall-template.S:82
#2  0xb7626eff in __sleep (seconds=0) at ../sysdeps/unix/sysv/linux/sleep.c:138
#3  0x080487c8 in crash_handler (sig=11) at hdtrap.c:73
#4  <signal handler called>
#5  0x08048930 in calc_md5 (data=0x8048a00 "testing data-xxxxxxx", nLen=20,
    md5=0x0) at md5.c:7
#6  0x08048698 in get_file_id (filename=0x8048abe "filea", fileid=0xbf8b57c0)
    at hdtrap.c:27
#7  0x0804891d in main (argc=2, argv=0xbf8b5874) at hdtrap.c:115
(gdb) █

```

58

时光倒流

```
(gdb) frame 5
#5 0x08048930 in calc_md5 (data=0x8048a00 "testing data-xxxxxx", nLen=20,
    md5=0x0) at md5.c:7
7      md5[0] = A;
(gdb) info locals
A = 0
B = 134515200
C = -1218498352
D = -1217308488
(gdb) p md5
$1 = (unsigned int *) 0x0
```

59

理解错位

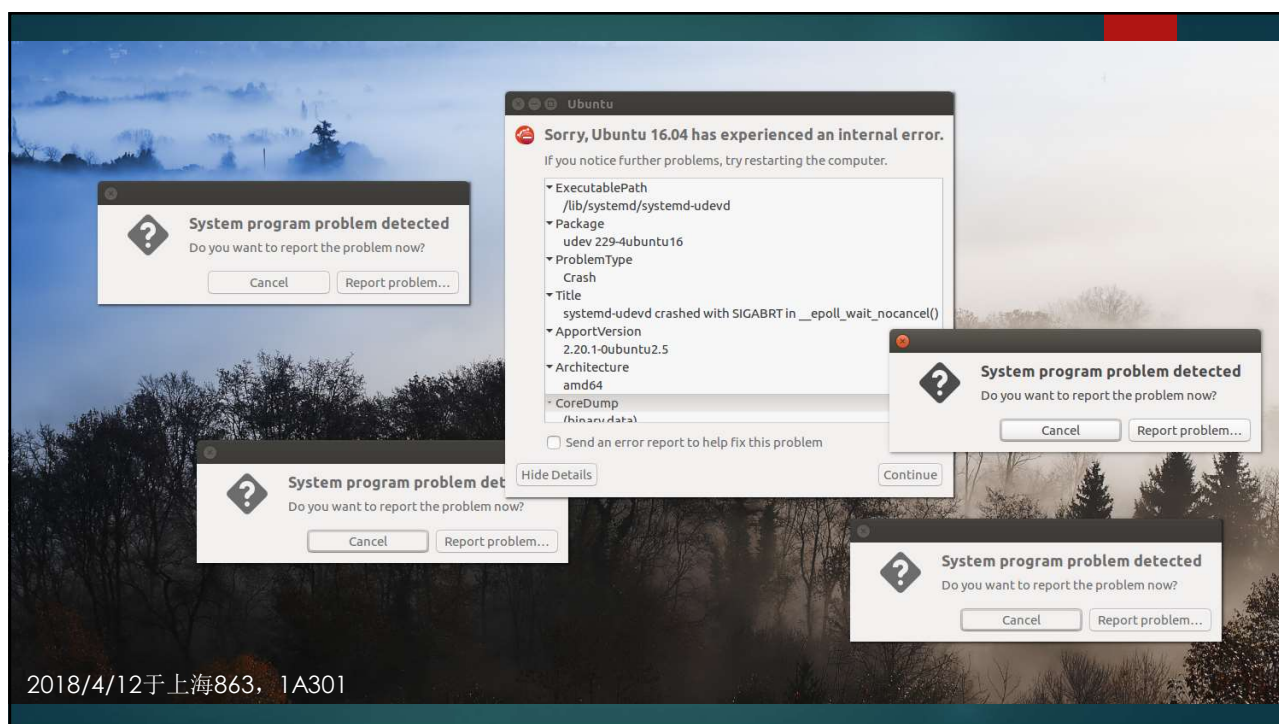
```
(gdb) x /32dx $esp
0xbf8b5768:    0xb77158b8    0xb75f30d0    0x08048a00    0x00000000
0xbf8b5778:    0xbf8b57a8    0x08048698    0x08048a00    0x00000014
0xbf8b5788:    0x00000000    0xbf8b57c0    0x0000001f    0x08048a00
0xbf8b5798:    0x00000014    0x00000000    0x00000000    0xb7713ff4
0xbf8b57a8:    0xbf8b57d8    0x0804891d    0x08048abe    0xbf8b57c0
0xbf8b57b8:    0xb7713ff4    0xb75a1e65    0xb7782230    0x00000000
0xbf8b57c8:    0x0804895b    0xb7713ff4    0x08048950    0x00000000
0xbf8b57d8:    0x00000000    0xb7588533    0x00000002    0xbf8b5874
```

```
(gdb) frame 7
#7 0x0804891d in main (argc=2, argv=0xbf8b5874) at hdtrap.c:115
115      get_file_id("filea", fileid);
(gdb) info locals
fileid = {3078103600, 0, 134515035, 3077652468}
(gdb) p &fileid
$2 = (unsigned int (*)[4]) 0xbf8b57c0
```

60




61



62

Apport

- ▶ <https://wiki.ubuntu.com/Apport>
- ▶ 自动介入
 - ▶ 崩溃
 - ▶ 其它可以自动检测到的错误情况，比如更新失败
- ▶ 收集错误信息
- ▶ 呈现界面



63

介入方式

- ▶ `/proc/sys/kernel/core_pattern`

```
gedu@gedu-VirtualBox:~$ cat /proc/sys/kernel/core_pattern
|/usr/share/apport/apport %p %s %c %P
```

- ▶ 上面的 | 代表用户态程序
- ▶ 对于Python程序，会安装/etc/python*/sitecustomize.py来调用apport

64

Apport脚本

- ▶ 用Python编写
- ▶ 作者Martin Pitt

```

Lister - [c:\dig\training\ltune\crash\apport]
File Edit Options Help
#!/usr/bin/python3

# Collect information about a crash and create a report in the directory
# specified by apport.Fileutils.report_dir.
# See https://wiki.ubuntu.com/Apport for details.
#
# Copyright (c) 2006 - 2016 Canonical Ltd.
# Author: Martin Pitt <martin.pitt@ubuntu.com>
#
# This program is free software; you can redistribute it and/or modify it
# under the terms of the GNU General Public License as published by the
# Free Software Foundation; either version 2 of the License, or (at your
# option) any later version. See http://www.gnu.org/copyleft/gpl.html for
# the full text of the license.

import sys, os, os.path, subprocess, time, traceback, pwd, io
import signal, inspect, grp, fcntl, socket, atexit, array, struct
import errno

import apport, apport.Fileutils

#####
#
# Functions
#
#####


def check_lock():
    '''Abort if another instance of apport is already running.

    This avoids bringing down the system to its knees if there is a series of
    crashes.'''

    # create a lock file
    lockfile = os.path.join(apport.Fileutils.report_dir, '.lock')
    try:
        fd = os.open(lockfile, os.O_WRONLY | os.O_CREAT | os.O_NOFOLLOW)
    except OSError as e:
        error_log('cannot create lock file (uid %i): %s' % (os.getuid(), str(e)))
        sys.exit(1)

```

65



MARTIN PITT
ADDICTED TO FREE SOFTWARE DEVELOPMENT


g+ ↻ ✉

Blog Documents FOSS Projects

The alphabet and pitti end here: Last day at Canonical

I've had the pleasure of working on Ubuntu for 12½ years now, and during that time used up an entire Latin alphabet of release names! (Well, A and C are still free, but we used H and W twice, so on average.. ☺) This has for sure been the most exciting time in my life with tons of good memories! Very few highlights:

- Getting some spam mail from a South African multi-millionaire about a GREAT OPPORTUNITY
- Joining #warthogs (my first IRC experience) and collecting my first bounties for "derooting" Debian (i. e. drop privileges from root daemons and suid binaries)
- Getting invited to Oxford to meet a bunch of people for which I had absolutely zero proof of existence, and tossing myself into debts for buying a laptop for that occasion
- Once being there, looking into my fellows' stern and serious faces and being amazed by their professionalism:



- The excitement and hype around going public with [Warty Warthogs Beta](#)

66

产生界面

- ▶ 在Gnome中， update-notifier进程会监视/var/crash目录

```
gedu@gedu-VirtualBox:~$ ps -A | grep update
2542 ?      00:00:00 update-notifier
```

- ▶ 一旦目录有变化，它会调用/usr/share/apport/apport-checkreports
- ▶ 如果有新报告，则调用/usr/share/apport/apport-gtk呈现界面

67

问题类型

Bug

- The user has run ubuntu-bug manually to report an issue.

Crash

- Apport automatically detected a program crash.

Hang

- Apport automatically detected a program has stopped responding.

Package

- Apport automatically detected a program in a package failed to install/upgrade correctly.

KernelOops

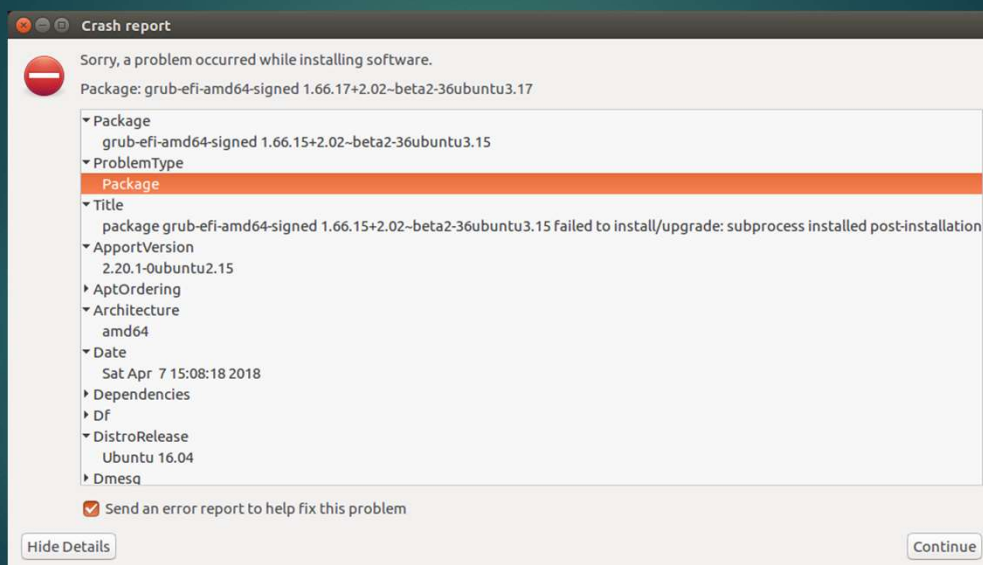
- Apport detected that the kernel encountered a situation that should not have arisen (non fatal).

KernelCrash

- Apport detected that the kernel crashed (on a previous boot).

68

包类型的故障



69

报告格式

```
DistroRelease: Ubuntu 12.04
ExecutablePath: /usr/bin/gcalctool
Package: gcalctool 5.8.24-0ubuntu2
ProcCmdline: gcalctool
ProcEnviron:
  SHELL=/bin/bash
  PATH=/usr/sbin:/usr/bin:/sbin:/bin:/usr/bin/X11:/usr/games
  LANG=de_DE.UTF-8
StackTrace:
[...]
#0 0x00002ae577bb37bf in poll () from /lib/libc.so.6
No symbol table info available.
#1 0x00002ae57786991e in g_main_context_check () from /usr/lib64/libglib-2.0.so.0
No symbol table info available.
[...]
CoreDump: base64
eJzsXQmcFMXV7+XGA0dBREVoDxSPXQYEB...
```

70

实例

```
gedu@gedu-VirtualBox:~$ cat /usr/share/apport/package-hooks/udev.py
"""apport package hook for udev
```

```
(c) 2009 Canonical Ltd.
Author: Martin Pitt <martin.pitt@ubuntu.com>
"""
```

```
import os
import apport.hookutils
```

```
def add_info(report):
    apport.hookutils.attach_hardware(report)
```

```
    user_rules = []
    for f in os.listdir('/etc/udev/rules.d'):
        if not f.startswith('70-persistent-') and f != 'README':
            user_rules.append(f)
```

```
    if user_rules:
        report['CustomUdevRuleFiles'] = ' '.join(user_rules)
```

73


可以提交互性问题

► 回答Yes和No

```
def add_info(report, ui):
    attach_file = False

    if ui and ui.yesno('Attach
some file?') == True:
        attach_file = True

    if attach_file == True:
        # user wants file to be
        attached
```



74

或者从列表选择

```
def add_info(report, ui):

    attach_file = False

    days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday' ]

    response = ui.choice('What day is it?', days)

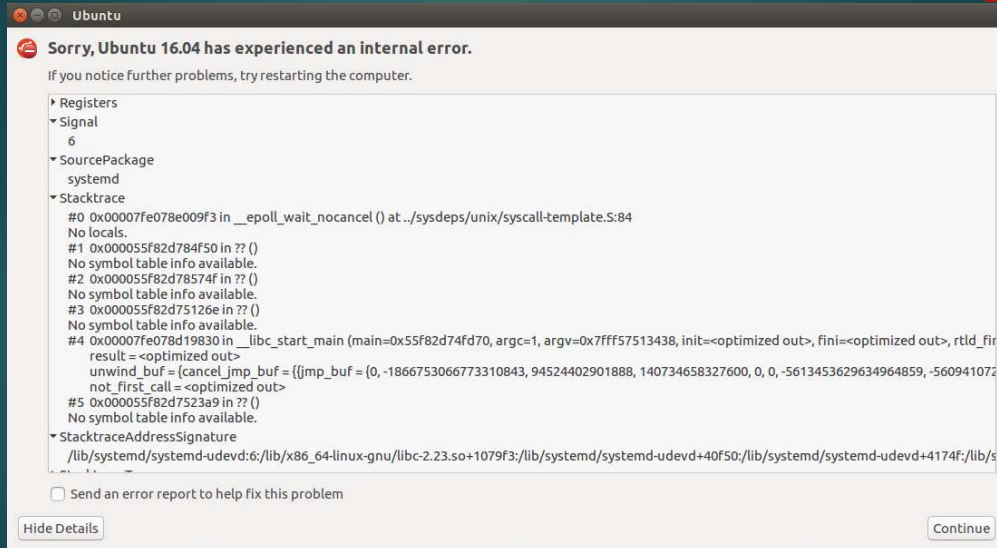
    if response == None:
        # user cancelled
        pass
    elif 0 in response:
        # 'Monday' selected
    elif 1 in response:
        # 'Tuesday' selected
    ...
    elif 6 in response:
        # 'Sunday' selected
```

75

```
# map crash database names to CrashDatabase implementations and URLs
default = 'ubuntu'
def get_oem_project():
    """Determine OEM project name from Distribution Channel Descriptor
    Return None if it cannot be determined or does not exist.
    """
    try:
        dcd = open('/var/lib/ubuntu_dist_channel').read()
        if dcd.startswith('canonical-oem-'):
            return dcd.split('-')[2]
    except IOError:
        return None
databases = {
    'ubuntu': {
        'impl': 'launchpad',
        'bug_pattern_url': 'http://people.canonical.com/~ubuntu-archive/bugpatterns/bugpatterns.xml',
        'dupdb_url': 'http://people.canonical.com/~ubuntu-archive/apport-duplicates',
        'distro': 'ubuntu',
        'problem_types': ['Bug', 'Package'],
        'escalation_tag': 'bugpattern-needed',
        'escalated_tag': 'bugpattern-written',
    },
    'canonical-oem': {
        'impl': 'launchpad',
```

76

案例: udevd进程崩溃



77

apport-retrace

- ▶ 需要安装
- ▶ `apport-retrace -g CRASHFILE.crash`

```
gedu@gedu-VirtualBox:~/Labs/crashes$ apport-retrace -g _lib_systemd_systemd-udev9.0.crash
GNU gdb (Ubuntu 7.11.1-0ubuntu1-16.04) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.  Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
Reading symbols from /lib/systemd/systemd-udev...(no debugging symbols found)...done.
[New LWP 271]
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Core was generated by "/lib/systemd/systemd-udev".
Program terminated with signal SIGABRT, Aborted.
#0  0x00007fe078e009f3 in __epoll_wait_nocancel ()
    at ../sysdeps/unix/syscall-template.S:84
#1  0x00007fe078e009f3 in __epoll_wait_nocancel ()
    at ../sysdeps/unix/syscall-template.S:84
(gdb) bt
#0  0x00007fe078e009f3 in __epoll_wait_nocancel ()
    at ../sysdeps/unix/syscall-template.S:84
#1  0x000055f82d784f50 in ?? ()
#2  0x000055f82d78574f in ?? ()
#3  0x000055f82d75126e in ?? ()
#4  0x00007fe078d19830 in _libc_start_main (main=0x55f82d74fd70, argc=1,
    argv=0x7ffff5713438, init=<optimized out>, fini=<optimized out>,
    rtd_fini=<optimized out>, stack_end=0x7ffff5713428)
    at ../csu/libc-start.c:291
#5  0x000055f82d7523a9 in ?? ()
(gdb)
```

78

解包

- ▶ `apport-unpack systemGeneratedCrashReportPath.crash yourNewUnpackDirectoryHere`
- ▶ `cd yourNewUnpackDirectoryHere/`
- ▶ `gdb `cat ExecutablePath` CoreDump`

79

SIGABRT

- ▶ `#define SIGIOT 6 /* IOT trap (4.2 BSD). */`
- ▶ `#define SIGABRT SIGIOT /* Abort (ANSI). */`

80


```

C:\Windows\System32\cmd.exe - adb shell
at android.os.HandlerThread.run(HandlerThread.java:61)

"RxCachedThreadScheduler-15" daemon prio=5 tid=15 Waiting
  group="main" sCount=1 dsCount=0 obj=0x12c10670 self=0xc92b9400
  sysTid=10421 nice=0 cgrp=bg_non_interactive sched=0/0 handle=0xd4f96920
  state=S schedstat=( 30387608 25068802 93 ) utm=3 stm=2 core=0 HZ=100
  stack=0xd4e94000-0xd4e96000 stackSize=1038KB
  held mutexes=
  at java.lang.Object.wait!(Native method)
  - waiting on <0x0felbaa8> (a java.lang.Object)
  at java.lang.Thread.parkFor$(Thread.java:2127)
  - locked <0x0felbaa8> (a java.lang.Object)
  at sun.misc.Unsafe.park(Unsafe.java:325)
  at java.util.concurrent.locks.LockSupport.park(LockSupport.java:161)
  at java.util.concurrent.locks.AbstractQueuedSynchronizer$ConditionObject.await(AbstractQueuedSynchronizer.java:2035)
  at java.util.concurrent.ScheduledThreadPoolExecutor$DelayedWorkQueue.take(ScheduledThreadPoolExecutor.java:1091)
  at java.util.concurrent.ScheduledThreadPoolExecutor$DelayedWorkQueue.take(ScheduledThreadPoolExecutor.java:1084)
  at java.util.concurrent.ThreadPoolExecutor.getTask(ThreadPoolExecutor.java:1058)
  at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1118)
  at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:607)
  at java.lang.Thread.run(Thread.java:761)

"amapLocManagerThread" prio=5 tid=16 Native
  group="main" sCount=1 dsCount=0 obj=0x12c0f420 self=0xc92b9e00
  sysTid=10423 nice=0 cgrp=bg_non_interactive sched=0/0 handle=0xd4e91920
  state=S schedstat=( 104352291 47637446 378 ) utm=8 stm=2 core=1 HZ=100
  stack=0xd4d8f000-0xd4d91000 stackSize=1038KB
  held mutexes=
  kernel: __switch_to+0x70/0x7c
  kernel: Sys epoll wait+0x29c/0x328

```

81

TombStones

- ▶ 默认在/data/tombstones文件夹
- ▶ 文件名为“tombstone_00”，“tombstone_10”
- ▶ 需要root权限才可以访问（糟糕！）

82

```

TIME=1315550899000
FINGERPRINT=samsung/GT-S5830/GT-S5830:2.3.4/GINGERBREAD/ZCKPB:user/release-keys
HARDWARE=gt-s5830
UNKNOWN=unknown
RADIO=unknown
BOARD=GT-S5830
versionCode=17
PRODUCT=GT-S5830
versionName=2.4.3
DISPLAY=GINGERBREAD.ZCKPB
USER=se.infra
HOST=SEP-76
DEVICE=GT-S5830
TAGS=release-keys
MODEL=GT-S5830
BOOTLOADER=unknown
CPU_ABI=armeabi
CPU_ABI2=unknown
ID=GINGERBREAD
SERIAL=unknown
MANUFACTURER=samsung
BRAND=samsung
TYPE=user

```

83

```

java.lang.NoClassDefFoundError: com.baidu.mapapi.BMapManager
    at com.hiad365.zyh.ZYHApplication.initEngineManager(ZYHApplication.java:80)
    at com.hiad365.zyh.ZYHApplication.onCreate(ZYHApplication.java:70)
    at android.app.Instrumentation.callApplicationOnCreate(Instrumentation.java:969)
    at android.app.ActivityThread.handleBindApplication(ActivityThread.java:3276)
    at android.app.ActivityThread.access$2200(ActivityThread.java:117)
    at android.app.ActivityThread$H.handleMessage(ActivityThread.java:973)
    at android.os.Handler.dispatchMessage(Handler.java:99)
    at android.os.Looper.loop(Looper.java:130)
    at android.app.ActivityThread.main(ActivityThread.java:3687)
    at java.lang.reflect.Method.invokeNative(Native Method)
    at java.lang.reflect.Method.invoke(Method.java:507)
    at com.android.internal.os.ZygoteInit$MethodAndArgsCaller.run(ZygoteInit.java:867)
    at com.android.internal.os.ZygoteInit.main(ZygoteInit.java:625)
    at dalvik.system.NativeStart.main(Native Method)

```

84

```
java.lang.ClassNotFoundException: com.baidu.mapapi.BMapManager in loader
dalvik.system.PathClassLoader[/data/app/com.hiad365.zyh-1.apk]
  at dalvik.system.PathClassLoader.findClass(PathClassLoader.java:240)
  at java.lang.ClassLoader.loadClass(ClassLoader.java:551)
  at java.lang.ClassLoader.loadClass(ClassLoader.java:511)
  at com.hiad365.zyh.ZYHApplication.initEngineManager(ZYHApplication.java:80)
  at com.hiad365.zyh.ZYHApplication.onCreate(ZYHApplication.java:70)
  at android.app.Instrumentation.callApplicationOnCreate(Instrumentation.java:969)
  at android.app.ActivityThread.handleBindApplication(ActivityThread.java:3276)
  at android.app.ActivityThread.access$2200(ActivityThread.java:117)
  at android.app.ActivityThread$H.handleMessage(ActivityThread.java:973)
  at android.os.Handler.dispatchMessage(Handler.java:99)
  at android.os.Looper.loop(Looper.java:130)
  at android.app.ActivityThread.main(ActivityThread.java:3687)
  at java.lang.reflect.Method.invokeNative(Native Method)
  at java.lang.reflect.Method.invoke(Method.java:507)
  at com.android.internal.os.ZygoteInit$MethodAndArgsCaller.run(ZygoteInit.java:867)
  at com.android.internal.os.ZygoteInit.main(ZygoteInit.java:625)
  at dalvik.system.NativeStart.main(Native Method)
```

85

切问而近思

欢迎关注格友公众号



86