

# LINUX系统高级调试和 优化

## --事件追踪

张银奎

2017/8/12

1

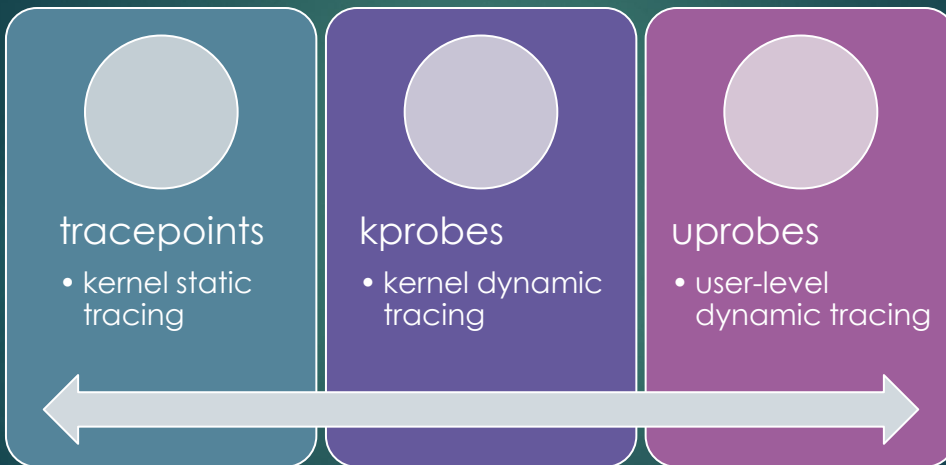
## Gregg的追踪机制时间表

- 1990's: Static tracers, prototype dynamic tracers
  - 2004: Linux kprobes (2.6.9)
    - Dynamic kernel tracing, difficult interface
  - 2005: Solaris DTrace (s10)
    - Static & dynamic tracing, user & kernel level, production ready, easy to use, far better than anything of the time, and, marketed
  - 2008: Linux ftrace (2.6.27)
  - 2009: Linux perf (2.6.31)
  - 2009: Linux tracepoints (2.6.32)
    - Static kernel tracing
  - 2010-2014: ftrace & perf\_events enhancements
  - 2014: eBPF patches
    - Fast (JIT'd) in kernel aggregations and programs

[http://www.brendangregg.com/Slides/LISA2014\\_LinuxPerfAnalysisNewTools.pdf](http://www.brendangregg.com/Slides/LISA2014_LinuxPerfAnalysisNewTools.pdf)

2

## 三类追踪信息源



3

c:\bench\linux-4.4.14\Documentation\trace\\*.\*

Name	Ext	Size	Date	Attr
↑[...]		<DIR>	2016/06/25 01:18	----
↑[postprocess]		<DIR>	2016/06/25 01:18	----
coresight	txt	15,011	2016/06/25 01:18	-a--
events	txt	17,546	2016/06/25 01:18	-a--
events-kmem	txt	5,389	2016/06/25 01:18	-a--
events-nmi	txt	1,561	2016/06/25 01:18	-a--
events-power	txt	3,494	2016/06/25 01:18	-a--
ftrace	txt	109,566	2016/06/25 01:18	-a--
ftrace-design	txt	14,383	2016/06/25 01:18	-a--
function-graph-fold	vim	1,355	2016/06/25 01:18	-a--
intel_th	txt	3,665	2016/06/25 01:18	-a--
kprobetrace	txt	7,082	2016/06/25 01:18	-a--
mmiotrace	txt	6,804	2016/06/25 01:18	-a--
ring-buffer-design	txt	30,648	2016/06/25 01:18	-a--
stm	txt	3,941	2016/06/25 01:18	-a--
tracepoint-analysis	txt	12,243	2016/06/25 01:18	-a--
tracepoints	txt	5,227	2016/06/25 01:18	-a--
uprobetracer	txt	6,447	2016/06/25 01:18	-a--

4

# ftrace

- ▶ Linux内核中的踪迹工具
  - ▶ 最初在 2.6.27 中出现
- ▶ 最初只能产生函数调用的踪迹
- ▶ 目前已发展为框架，支持扩展
  - ▶ **Power tracer**: 记录系统电源管理相关的信息
  - ▶ **Hardware branch tracer**: 利用处理器的分支跟踪能力，实现硬件级别的指令跳转记录
  - ▶ **Schedule switch tracer**: 跟踪进程调度情况
  - ▶ **Initcall tracer**: 记录系统在 boot 阶段所调用的 init call

5

## ftrace - Function Tracer =====

Copyright 2008 Red Hat Inc.

Author: Steven Rostedt <ststedt@redhat.com>

License: The GNU Free Documentation License, Version 1.2  
(dual licensed under the GPL v2)

Reviewers: Elias Oltmanns, Randy Dunlap, Andrew Morton,  
John Kacur, and David Teigland.

Written for: 2.6.28-rc2


Updated for: 3.10

### Introduction -----

Ftrace is an internal tracer designed to help out developers and designers of systems to find what is going on inside the kernel. It can be used for debugging or analyzing latencies and performance issues that take place outside of user-space.

Although ftrace is typically considered the function tracer, it is really a frame work of several assorted tracing utilities. There's latency tracing to examine what occurs between interrupts disabled and enabled, as well as for preemption and from a time

6



<https://chemnitzer.linux-tage.de/2012/vortraege/media/videos/v4sa1500.mp4>

7

## 实现: kernel/trace

blktrace.c	ftrace.c	Kconfig
bpf_trace.c	power-traces.c	ring_buffer.c
Makefile	rpm-traces.c	trace.c
ring_buffer_benchmark.c	trace_benchmark.c	trace_benchmark.h
trace.h	trace_clock.c	trace_entries.h
trace_branch.c	trace_events_filter.c	trace_events_filter_test.h
trace_events.c	trace_event_perf.c	trace_export.c
trace_events_trigger.c	trace_functions_graph.c	trace_irqsoff.c
trace_functions.c	trace_kprobe.c	trace_mmioTRACE.c
trace_kdb.c	trace_output.c	trace_output.h
trace_nop.c	trace_probe.c	trace_probe.h
trace_printk.c	trace_sched_wakeup.c	trace_selftest.c
trace_sched_switch.c	trace_seq.c	trace_stack.c
trace_selftest_dynamic.c	trace_stat.h	trace_syscalls.c
trace_stat.c		
trace_uprobe.c		

44 个文件

1,109,715 字节

8

## VFS接口

Ftrace uses the debugfs file system to hold the control files as well as the files to display output.

When debugfs is configured into the kernel (which selecting any ftrace option will do) the directory /sys/kernel/debug will be created. To mount this directory, you can add to your /etc/fstab file:

```
debugfs    /sys/kernel/debug    debugfs defaults    0    0
```

Or you can mount it at run time with:

```
mount -t debugfs nodev /sys/kernel/debug
```

9

## 文件接口

```
root@gewubox:/sys/kernel/debug/tracing# ls
available_events          max_graph_depth          stack_trace_filter
available_filter_functions options                   trace
available_tracers         per_cpu                   trace_clock
buffer_size_kb            printk_formats            trace_marker
buffer_total_size_kb      README                   trace_options
current_tracer            saved_cmdlines            trace_pipe
dyn_ftrace_total_info     set_event                 trace_stat
enabled_functions         set_ftrace_filter         tracing_cpumask
events                   set_ftrace_notrace        tracing_max_latency
free_buffer               set_ftrace_pid            tracing_on
function_profile_enabled  set_graph_function        tracing_thresh
instances                 snapshot                   uprobe_events
kprobe_events             stack_max_size            uprobe_profile
kprobe_profile            stack_trace
```

10


## 使用ftrace

- ▶ 选择 tracer
- ▶ 启用
- ▶ 执行目标程序，促发产生踪迹信息
- ▶ 关闭 ftrace
- ▶ 查看 trace 文件

11

## 追踪器(tracer)

- ▶ 将追踪源的原始信息搜集起来，计算加工，并以标准接口提供给前端



12

## 数据结构

```
(gdb) pt function_trace
type = struct tracer {
  const char *name;
  int (*init)(struct trace_array *);
  void (*reset)(struct trace_array *);
  void (*start)(struct trace_array *);
  void (*stop)(struct trace_array *);
  void (*open)(struct trace_iterator *);
  void (*pipe_open)(struct trace_iterator *);
  void (*wait_pipe)(struct trace_iterator *);
  void (*close)(struct trace_iterator *);
  void (*pipe_close)(struct trace_iterator *);
  ssize_t (*read)(struct trace_iterator *, struct file *, char *, size_t, loff_t *);
  ssize_t (*splice_read)(struct trace_iterator *, struct file *, loff_t *,
  struct pipe_inode_info *, size_t, unsigned int);
  void (*print_header)(struct seq_file *);
  enum print_line_t (*print_line)(struct trace_iterator *);
  int (*set_flag)(u32, u32, int);
  int (*flag_changed)(struct tracer *, u32, int);
  struct tracer *next;
  struct tracer_flags *flags;
  bool print_max;
  bool enabled;
  bool use_max_tr;
}
```

13

## 格物

```
(gdb) p function_trace
$1 = {name = 0xc18752a4 "function", init = 0xc10ffb80 <function_trace_init>,
  reset = 0xc10ffb40 <function_trace_reset>,
  start = 0xc10ffb30 <function_trace_start>, stop = 0, open = 0,
  pipe_open = 0, wait_pipe = 0xc10fbc20 <poll_wait_pipe>, close = 0,
  pipe_close = 0, read = 0, splice_read = 0, print_header = 0, print_line = 0,
  set_flag = 0xc10ffc90 <func_set_flag>, flag_changed = 0, next = 0xc19e11e0,
  flags = 0xc197224c, print_max = false, enabled = false, use_max_tr = false}
```

14



## 观察可用的tracer

```
root@gewubox:/sys/kernel/debug/tracing# cat available_tracers
blk mmiotrace function_graph wakeup_rt wakeup function nop
```

15

## 选用tracer

- ▶ # sudo su
- ▶ # cd /sys/kernel/debug/tracing
- ▶ # echo function > current\_tracer
- ▶ #cat trace
  
- ▶ 快速禁止或者启用
- ▶ # echo 0 > tracing\_on : quick way to disable tracing
- ▶ # echo 1 > tracing\_on : quick way to re-enable tracing
  
- ▶ 停用
- ▶ # echo nop > current\_tracer

16



## 过滤

- ▶ tracing\_cpumask:
  - ▶ This is a mask that lets the user only trace on specified CPUs. The format is a hex string representing the CPUs.
- ▶ set\_ftrace\_notrace:
  - ▶ This has an effect opposite to that of
- ▶ set\_ftrace\_filter.
  - ▶ Any function that is added here will not be traced. If a function exists in both set\_ftrace\_filter and set\_ftrace\_notrace, the function will not be traced.
- ▶ set\_ftrace\_pid:
  - ▶ Have the function tracer only trace a single thread.
- ▶ set\_event\_pid:
  - ▶ Have the events only trace a task with a PID listed in this file.
  - ▶ Note, sched\_switch and sched\_wake\_up will also trace events listed in this file.

17

## trace\_marker

```
void trace_write(const char *fmt, ...)
{
    va_list ap;
    char buf[256];
    int n;

    if (trace_fd < 0)
        return;

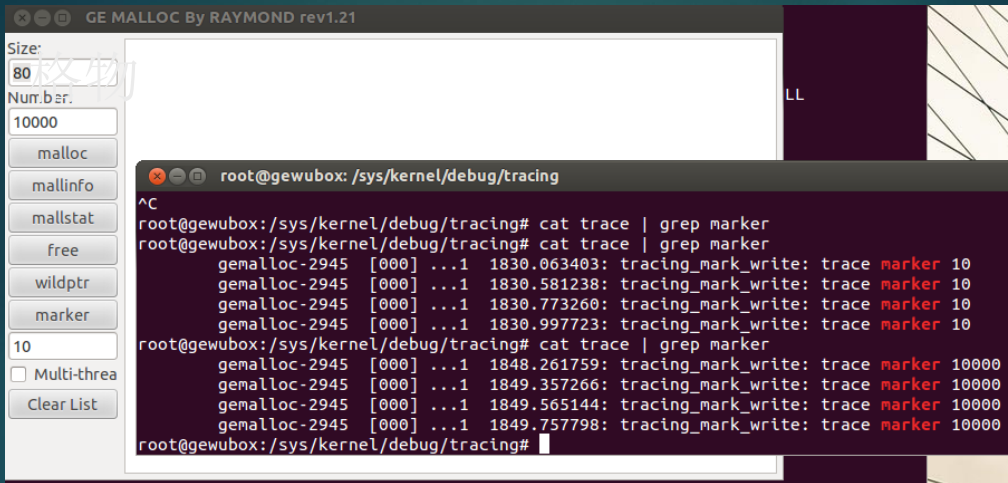
    va_start(ap, fmt);
    n = vsnprintf(buf, 256, fmt, ap);
    va_end(ap);

    write(trace_fd, buf, n);
}

start:

    trace_fd = open("trace_marker", WR_ONLY);
```

18



```
# cd /home/ge/work/gemalloc
# sudo ./gemalloc
# echo function > current_tracer
# echo pid_of_gemalloc > set_fftrace_pid
# cat trace | grep marker
```

19

```
# tracer: function
#
# entries-in-buffer/entries-written: 140080/250280  #P:4
#
#          _-----=> irqsoff
#          / _-----=> need-resched
#          | / _---=> hardirq/softirq
#          || / _--=> preempt-depth
#          ||| / _delay
#
# TASK-PID CPU#  ||||  TIMESTAMP FUNCTION
#   ||   ||||   ||   |
bash-1977 [000] .... 17284.993652: sys_close <-system_call_fastpath
bash-1977 [000] .... 17284.993653: __close_fd <-sys_close
bash-1977 [000] .... 17284.993653: _raw_spin_lock <-__close_fd
sshd-1974 [003] .... 17284.993653: __srcu_read_unlock <-fsnotify
bash-1977 [000] .... 17284.993654: add_preempt_count <-_raw_spin_lock
bash-1977 [000] ...1 17284.993655: _raw_spin_unlock <-__close_fd
bash-1977 [000] ...1 17284.993656: sub_preempt_count <-_raw_spin_unlock
bash-1977 [000] .... 17284.993657: filp_close <-_close_fd
bash-1977 [000] .... 17284.993657: dnotify_flush <-filp_close
sshd-1974 [003] .... 17284.993658: sys_select <-system_call_fastpath
```

Note: all time values are in microseconds.

20

```

# tracer: irqsoff
#
# irqsoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 16 us, #4/4, CPU#0 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: swapper/0-0 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: run_timer_softirq
# => ended at: run_timer_softirq
#
#
#          _-----=> CPU#
#          / _-----=> irqsoff
#          | / _-----=> need-resched
#          || / _-----=> hardirq/softirq
#          ||| / _-----=> preempt-depth
#          |||| / _-----=> delay
# cmd pid ||||| time | caller
# \ / ||||| \ | /
<idle>-0 0d.s2 0us+: _raw_spin_lock_irq <-run_timer_softirq
<idle>-0 0dNs3 17us : _raw_spin_unlock_irq <-run_timer_softirq
<idle>-0 0dNs3 17us+: trace_hardirqs_on <-run_timer_softirq
<idle>-0 0dNs3 25us : <stack trace>
=> raw spin unlock irq

```

21

## trace-cmd: A front-end for Ftrace

trace-cmd [COMMAND] ...

- ▶ record - record a trace into a trace.dat file
- ▶ start - start tracing without recording into a file
- ▶ extract - extract a trace from the kernel
- ▶ stop - stop the kernel from recording trace data
- ▶ reset - disable all kernel tracing and clear the trace buffers
- ▶ report - read out the trace stored in a trace.dat file
- ▶ split - parse a trace.dat file into smaller file(s)
- ▶ listen - listen on a network socket for trace clients
- ▶ list - list the available events, plugins or options

22

```
trace-cmd start -p function
```

```
#0 function_trace_start (tr=0xc1b5f180) at kernel/trace/trace_functions.c:46
#1 0xc10f58d0 in rb_simple_write (filp=<optimized out>, ubuf=<optimized out>,
cnt=1, ppos=0xe9193f98) at kernel/trace/trace.c:5845
#2 0xc1178645 in vfs_write (file=0xe28919c0, buf=0x806389b "1", count=1,
pos=0xe9193f98) at fs/read_write.c:473
#3 0xc1178af7 in SYSC_write (count=1, buf=0x806389b "1", fd=<optimized out>)
at fs/read_write.c:522
#4 SyS_write (fd=3, buf=134625435, count=1) at fs/read_write.c:515
#5 <signal handler called>
#6 0xb77ad424 in ?? ()
#7 0xb7602533 in ?? ()
```

## plugin function

23

- Left click and drag left

# KernelShark

24

# 安装

```


root@gewubox:~# sudo apt-get install kernelshark
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  kernelshark
0 upgraded, 1 newly installed, 0 to remove and 358 not upgraded.
Need to get 263 kB of archives.
After this operation, 586 kB of additional disk space will be used.
Get:1 http://cn.archive.ubuntu.com/ubuntu/ precise/universe kernelshark i386 1.0.3-0ubuntu2 [263 kB]
Fetched 263 kB in 0s (730 kB/s)
Selecting previously unselected package kernelshark.
(Reading database ... 151009 files and directories currently installed.)
Unpacking kernelshark (from .../kernelshark_1.0.3-0ubuntu2_i386.deb) ...
Setting up kernelshark (1.0.3-0ubuntu2) ...

```

25

# 不完整版本


- ▶ 缺少Record和help菜单



26

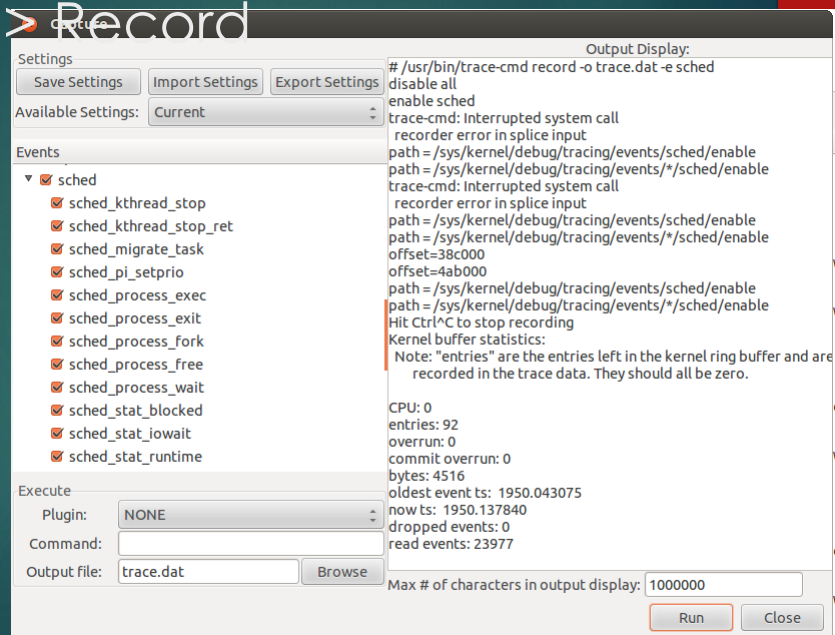
# 自己build

- ▶ `git clone git://git.kernel.org/pub/scm/linux/kernel/git/rostedt/trace-cmd.git`
- ▶ `sudo apt-get install swig`
- ▶ `sudo apt-get install libxml2-dev`
- ▶ `make gui`



27

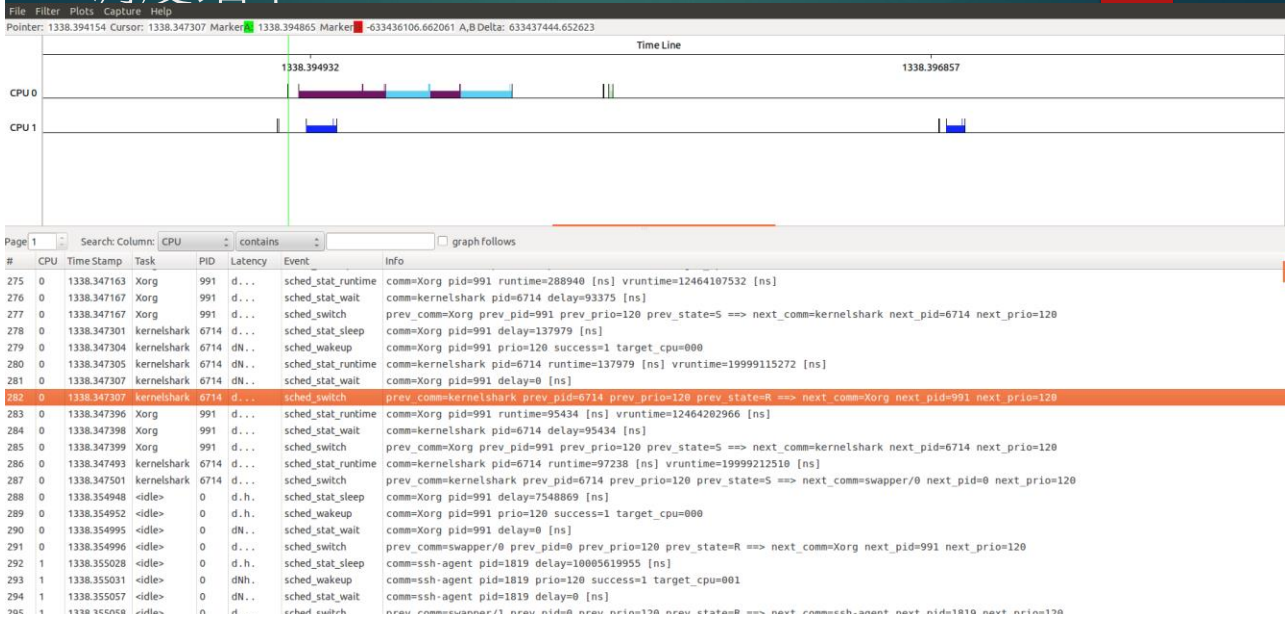
## Capture > Record



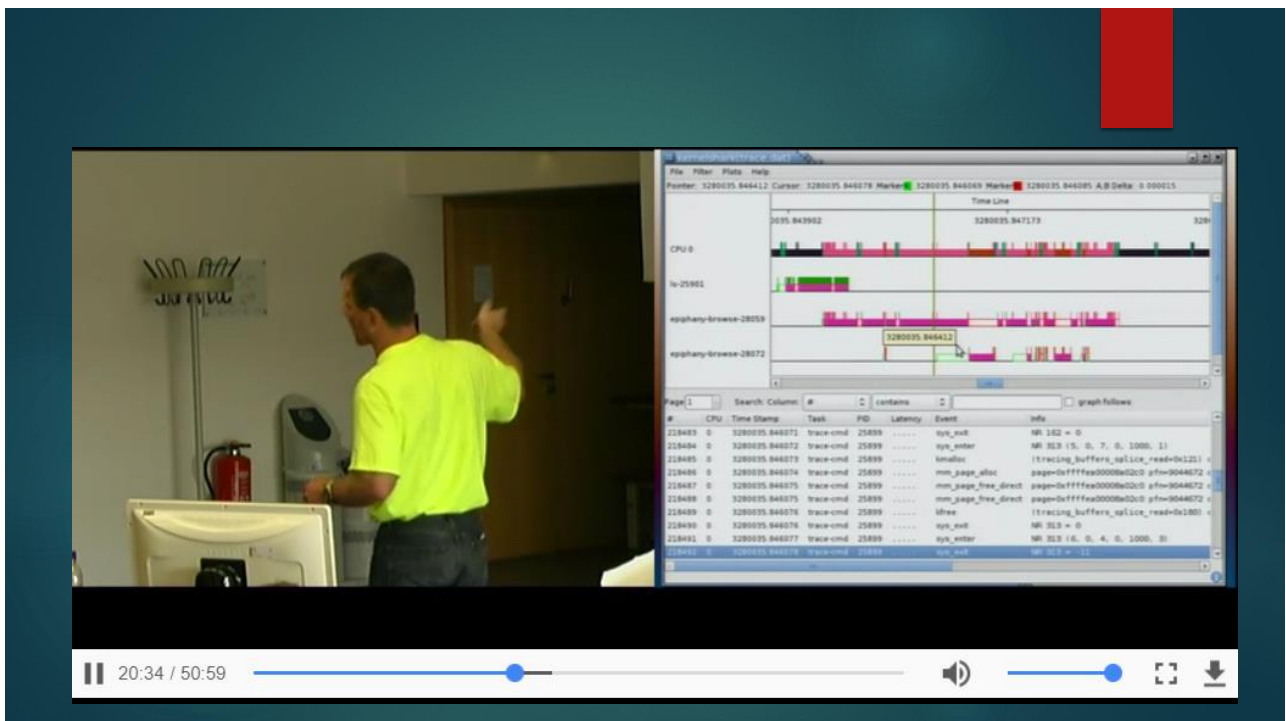
28



# 调度细节



29

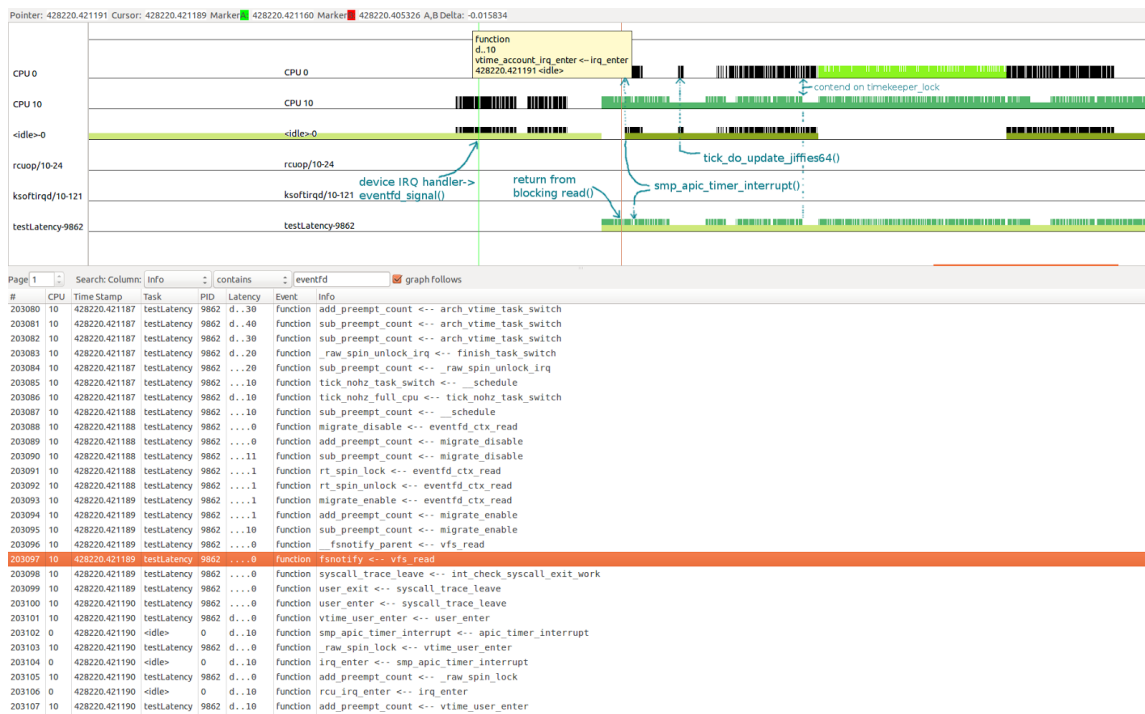


30





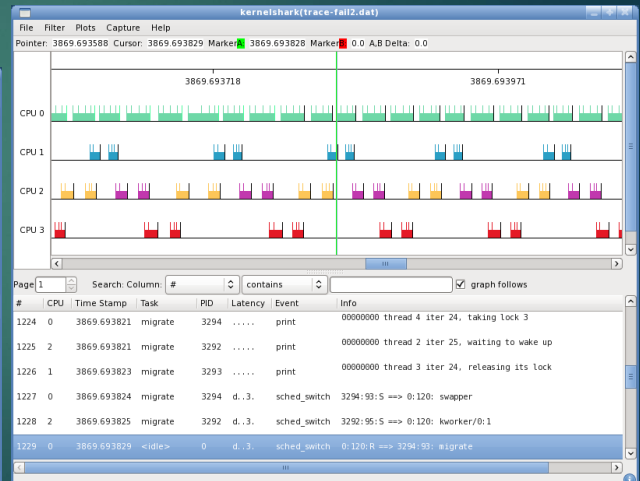
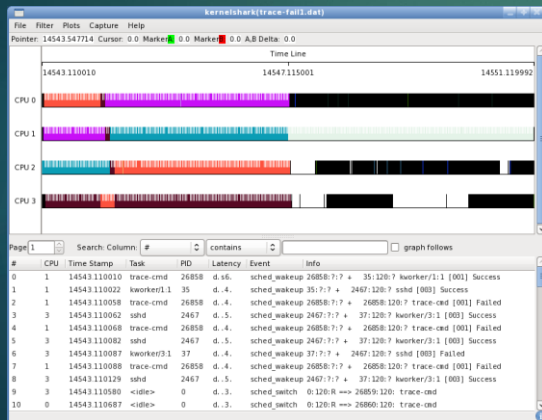
31



32

## Using KernelShark to analyze the real-time scheduler

- ▶ <https://lwn.net/Articles/425583/>
- ▶ Steven Rostedt的文章



33

## 资源

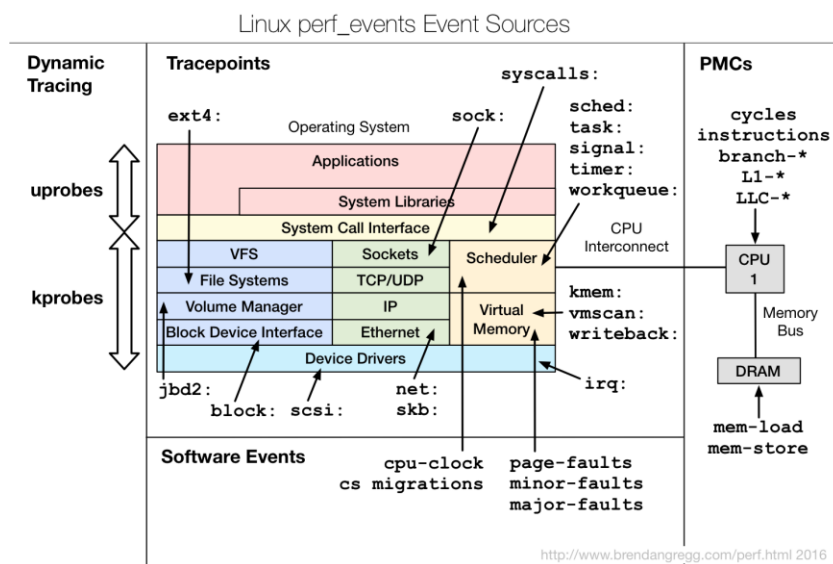
- ▶ LWN two-part article on [ftrace basics](#).
- ▶ Roasted's LWN.net [article](#) on how to use kernelshark.
- ▶ Documentation distributed with the Linux source code. Look in [Documentation/trace](#).
- ▶ Roasted's LWN [article on trace-cmd](#).

34

# Perf events

- ▶ perf Linux profiler
- ▶ Performance Counters for Linux (PCL), Linux perf events (LPE), or perf\_events
- ▶ Both ftrace and perf are core Linux tracing tools, included in the kernel source. Your system probably has ftrace already, and perf is often just a package add.
- ▶ 安装linux-tools-common
- ▶ 或者自己构建
  - ▶ Source code is in Linux: tools/perf

35



36

# 构建

- ▶ cd linux-3.12.2/kernel/tools/perf
- ▶ sudo apt-get install flex
- ▶ sudo apt-get install bison
- ▶ make

```
CC tests/hists_link.o
CC tests/python-use.o
CC tests/bp_signal.o
CC tests/bp_signal_ovr.o
CC tests/task-exit.o
CC tests/sw-clock.o
CC tests/perf-time-to-tsc.o
CC tests/code-reading.o
CC tests/sample-parsing.o
CC tests/parse-no-sample-id-all.o
CC arch/x86/util/header.o
CC arch/x86/util/tsc.o
CC util/symbol-minimal.o
CC tests/keep-tracking.o
AR libperf.a
SUBDIR /home/ge/work/linux-3.12.2/t
CC debugfs.o
AR liblk.a
SUBDIR /home/ge/work/linux-3.12.2/t
* new build flags or cross compiler
CC FPIC          event-parse.o
CC FPIC          trace-seq.o
CC FPIC          parse-filter.o
CC FPIC          parse-utils.o
CC FPIC          kbuffer-parse.o
BUILD STATIC LIB libtraceevent.a
LINK perf
GEN perf-archive
```

37

## perf [--version] [--help] COMMAND [ARGS]

The most commonly used perf commands are:

annotate	Read perf.data (created by perf record) and display annotated code
archive	Create archive with object files with build-ids found in perf.data file
bench	General framework for benchmark suites
buildid-cache	Manage build-id cache.
buildid-list	List the buildids in a perf.data file
diff	Read perf.data files and display the differential profile
evlist	List the event names in a perf.data file
inject	Filter to augment the events stream with additional information
kmem	Tool to trace/measure kernel memory(slab) properties
kvm	Tool to trace/measure kvm guest os
list	List all symbolic event types
lock	Analyze lock events
mem	Profile memory accesses
record	Run a command and record its profile into perf.data
report	Read perf.data (created by perf record) and display the profile
sched	Tool to trace/measure scheduler properties (latencies)
script	Read perf.data (created by perf record) and display trace output
stat	Run a command and gather performance counter statistics
test	Runs sanity tests.
timechart	Tool to visualize total system behavior during a workload
top	System profiling tool.
trace	strace inspired tool

38

```
ge@gewubox:~/work/linux-3.12.2/tools/perf$ ./perf list
```

List of pre-defined events (to be used in -e):

cpu-clock	[Software event]
task-clock	[Software event]
page-faults OR faults	[Software event]
context-switches OR cs	[Software event]
cpu-migrations OR migrations	[Software event]
minor-faults	[Software event]
major-faults	[Software event]
alignment-faults	[Software event]
emulation-faults	[Software event]
dummy	[Software event]
 rNNN	 [Raw hardware event descriptor]
cpu/t1=v1[, t2=v2, t3 ...]/modifier	[Raw hardware event descriptor]
(see 'man perf-list' on how to encode it)	
 mem:<addr>[:access]	 [Hardware breakpoint]

39

## perf record -e page-faults -a

```
ge@gewubox:/tmp/kshark$ sudo perf report
```

```
# =====
# captured on: Thu Aug 17 17:51:48 2017
# hostname : gewubox
# os release : 3.12.2
# perf version : 3.12.2
# arch : i686
# nrcpus online : 2
# nrcpus avail : 2
# cpudesc : Intel(R) Core(TM) i5-7200U CPU @ 2.50GHz
# cpuid : GenuineIntel,6,142,9
# total memory : 1544896 kB
# cmdline : /usr/bin/perf record -e page-faults -a
# event : name = page-faults, type = 1, config = 0x2, config1 = 0x0, config2 = 0x0, excl_usr = 0, excl_kern
# HEADER_CPU_TOPOLOGY info available, use -l to display
# pmu mappings: software = 1, tracepoint = 2, breakpoint = 5
# =====
#
```

40

续

```
# Samples: 1K of event 'page-faults'
# Event count (approx.): 115486
#
# Overhead      Command          Shared Object          Symbol
# .....
#
 82.58%          Xorg      libpixmap-1.so.0.30.2  [.] 0x0005f6eb
  4.16%          bash      bash                  [.] 0x00046040
  2.20%          bash      libc-2.15.so          [.] 0x000768db
  1.36%          sed       libc-2.15.so          [.] 0x000d9910
  1.23%          bash      ld-2.15.so            [.] 0x0000b9d9
  1.13%          sed       ld-2.15.so            [.] 0x00005048
  0.71%  gnome-terminal ld-2.15.so            [.] 0x00018907
  0.66%  gnome-terminal libc-2.15.so          [.] 0x00137331
  0.63%          sed       bash                  [.] 0x0008b0b0
  0.48%  gnome-terminal libcairo.so.2.11000.2 [.] 0x00042725
  0.37%          ls        libc-2.15.so          [.] 0x0007d260
  0.31%          Xorg      libc-2.15.so          [.] 0x00077341
  0.27%          metacity  libc-2.15.so          [.] 0x00074e22
  0.26%  gnome-terminal libgobject-2.0.so.0.3200.4 [.] 0x0000cbbc
  0.21%          sed       [kernel.kallsyms]     [k] 0xc132b138
  0.19%          uname      ld-2.15.so            [.] 0x0000e180
  0.18%          sed       libselinux.so.1       [.] 0x00015930
  0.17%          ls        ld-2.15.so            [.] 0x000112f0
  0.15%          ls        ld-2.15.so            [.] 0x00000467
```

41

```
root@gewubox:/tmp# perf record -e page-faults -a
^C[ perf record: Woken up 1 times to write data ]
[ perf record: Captured and wrote 0.469 MB perf.data (~20500 samples) ]
root@gewubox:/tmp# ls
perf.data      pulse-PKdhtXMr18n  ssh-HGRLEuW1735
keyring-UHcMhw pulse-2L9K88eMLGn7  pulse-sYomDRVlhJlw  unity_support_test.1
root@gewubox:/tmp# sudo perf report
# =====
# captured on: Fri Dec  1 21:50:47 2017
# hostname : gewubox
# os release : 3.12.2
# perf version : 3.12.2
# arch : i686
# nrcpus online : 1
# nrcpus avail : 1
# cpudesc : Intel(R) Core(TM) i3-2100 CPU @ 2.90GHz
# cpuid : GenuineIntel,6,142,9
# total memory : 1544952 kB
# cmdline : /usr/bin/perf record -e page-faults -a
# event : name = page-faults, type = 1, config = 0x2, config1 = 0x0, config2 = 0x0, excl_usr = 0, excl_kern = 0, excl_guest = 0
# HEADER_CPU_TOPOLOGY info available, use -I to display
# pmu mappings: software = 1, tracepoint = 2, breakpoint = 5
# =====
# Samples: 49 of event 'page-faults'
# Event count (approx.): 5533
#
# Overhead      Command          Shared Object          Symbol
# .....
#
 95.57%          Xorg      libpixmap-1.so.0.30.2  [.] 0x0005f6f5
  3.65%  gnome-terminal libc-2.15.so          [.] 0x00076b1e
  0.58%  gnome-terminal libcairo.so.2.11000.2 [.] 0x0000cfbe
  0.14%          perf      libc-2.15.so          [.] 0x000e3080
```

demo

使用perf统计缺页异常


/tmp# perf record -e page-faults -a

(CTRL + C 终止)

perf report

42





43

## Try it

```
# git clone https://github.com/brendangregg/FlameGraph # or download it from github
# cd FlameGraph
# perf record -F 99 -a -g -- sleep 60
# perf script | ./stackcollapse-perf.pl > out.perf-folded
# ./flamegraph.pl out.perf-folded > perf-kernel.svg
```

```
ge@gewubox:~/work/FlameGraph$ sudo perf record -F 99 -a -g -- sleep 60
[ perf record: Woken up 5 times to write data ]
[ perf record: Captured and wrote 1.759 MB perf.data (~76871 samples) ]
[vboxguest] with build id 90d523cfb7053db5f4024d1c38d2e9a2b7bfa50e not found, continuing without
ge@gewubox:~/work/FlameGraph$ perf script | ./stackcollapse-perf.pl > out.perf_folded
failed to open perf.data: Permission denied
ge@gewubox:~/work/FlameGraph$ sudo perf script | ./stackcollapse-perf.pl > out.perf_folded
Filtering for events of type: cpu-clock
ge@gewubox:~/work/FlameGraph$ ./flamegraph.pl out.perf_folded > perf-kernel.svg
ge@gewubox:~/work/FlameGraph$ firefox perf-kernel.svg
```

44



切问而近思

欢迎关注格友公众号

